

An integrated system for privacy-preserving, and auditable transactions on Hyperledger Fabric

MPhil Thesis Defence

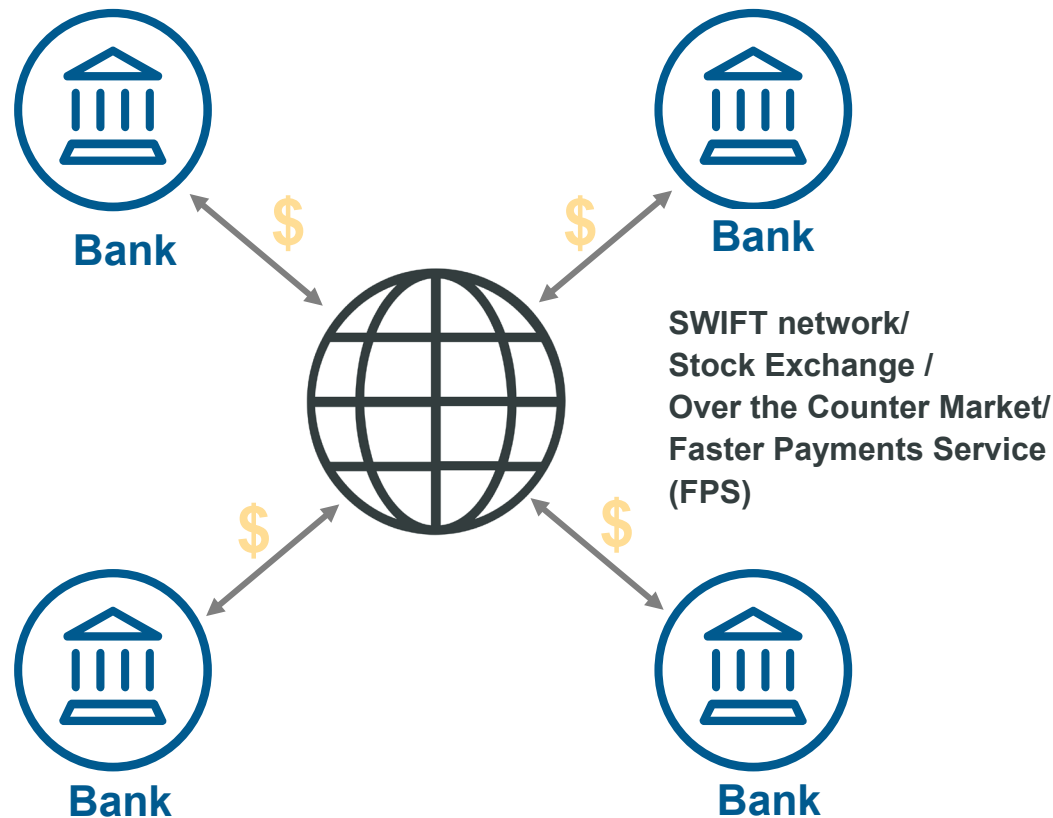
Presented by Martin Yip
Supervised by Prof. Dimitrios Papadopoulos

Introduction

Financial Service Industry

The financial service industry employs **centralized inter-bank transaction systems** [HKMA .01] to facilitates communication, transactions, and settlement among various banks through a central authority or platform.

Centralized Inter-Bank Transaction System:



[Benefit]

- Standardization
- Facilitate global trade
- Fraud detection

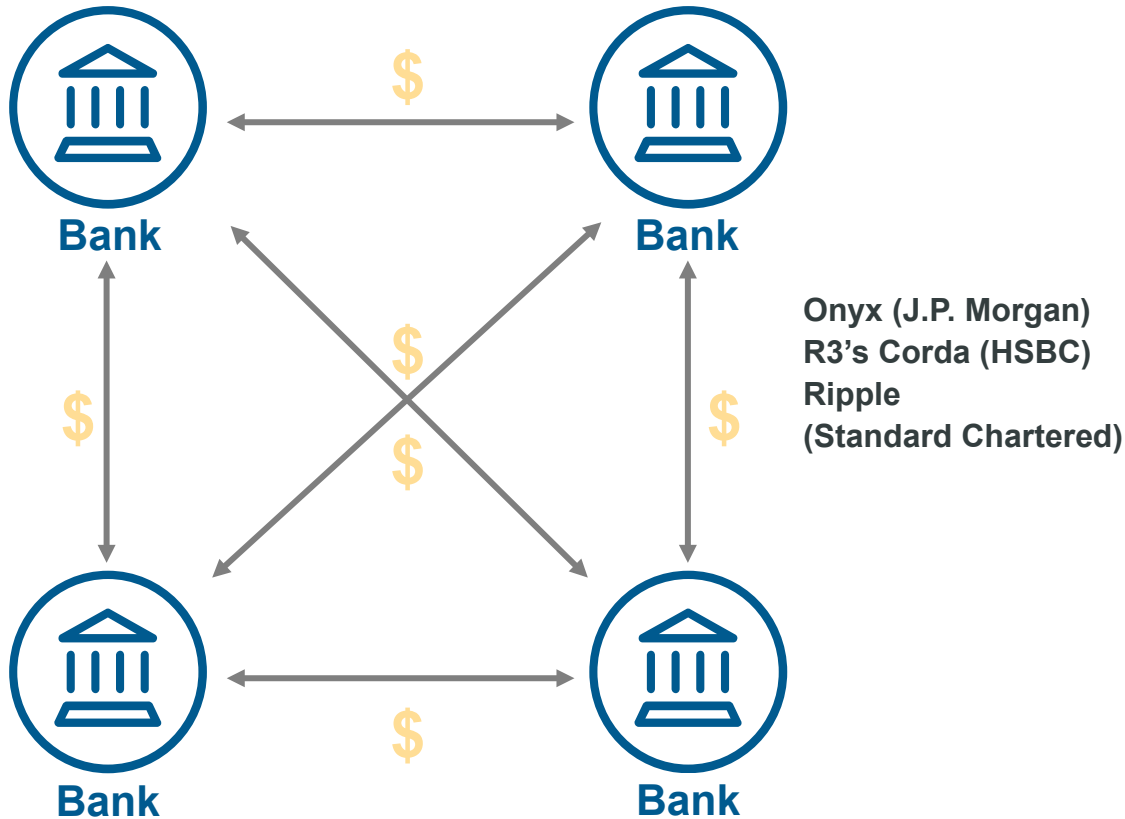
[Deficiency]

- Limited transparency
- Less flexibility
- Single point of failure

Financial Service Industry

The Hong Kong Monetary Authority (HKMA) has launched the [e-HKD Pilot Programme](#) to study the feasibility of retail central bank digital currencies as part of its “Fintech 2025” strategy [HKMA .23], including using [blockchain based platform](#) for transactions and settlement of tokenized assets.

Decentralized Inter-Bank Transaction System (Permissioned Blockchain):



[Benefit]

- System resilience
- Improved transparency
- Cost management

[Deficiency]

- Less scalable
- Adoption challenges
- Privacy concern


[Privacy Concern]

Banks concern that the transparency from collaborating will [reveal sensitive trading strategy information](#), potentially compromising their competitive advantage and [privacy](#).

Privacy-Preserving Transactions

Banks would use [privacy-preserving transactions](#) to hide the transaction amounts and sender/receivers' identities.

Clear transactions → Privacy-preserving transactions



Transfer transaction

Spending bank	SCB
Receiving bank	Goldman Sachs
Transfer amount	\$ 1,000,000



Transfer transaction

Spending bank	
Receiving bank	
Transfer amount	

Any problem if all of the transaction information is obscured?

Auditing & Privacy



Regardless of which transaction system the banks use, **financial audits** are required to be conducted regularly, which cover various aspects of a bank's operations, including its **balance sheet** to fulfill regulatory requirements. However, it will be challenging to audit ledgers with **privacy-preserving transactions**.

ID	Date	Asset	From	To	Amount
23	30 Nov 2023	\$	HSBC	SCB	1,000,000
24	30 Nov 2023	\$	BOC	Goldman Sachs	2,000,000
25	30 Nov 2023	\$	UBS	Citibank	1,000,000



[Privacy-Preserving Transactions]

ID	Date	Asset	From	To	Amount



Auditors will examine the ledger to verify to following financial invariants:

- ✓ Consent to transfer
- ✓ Has assets to transfer
- ✓ Assets neither created nor destroyed

How can the auditor verify the transactions now?

Related Works



Privacy-preserving transactions with auditability is a significant area of research in the blockchain and cryptography domains. Here are some research papers that focus on auditable privacy-preserving transactions in blockchain:

Paper	Benefit	Deficiency
zkLedger [Narula et al. 18]	a privacy-preserving, auditable distributed ledger system, allowing for regulatory compliance and auditability	has not been implemented on a blockchain platform.
MINILEDGER [Chatzigiannis et al. 21]	a private and auditable payment system with storage costs independent of the number of transactions.	has not been implemented on a blockchain platform.
FabZK [Kang et al. 19]	a privacy-preserving framework for Hyperledger Fabric, a permissioned blockchain	does not support a rich set of auditing primitives.
smartFHE [Solomon et al. 23]	a framework to support private smart contracts using fully homomorphic encryption	does not support a rich set of auditing primitives.
Zether [Bünz et al. 20]	a confidential payment mechanism compatible with Ethereum and other smart contract platforms.	does not support a rich set of auditing primitives.
Solidus [Cecchetti et al. 14]	a protocol for confidential transactions on public blockchains	does not support private auditing
Zerocash [Ben-Sasson et al. 14]	a protocol that extends Bitcoin enabling anonymous and non-traceable transactions	does not support private auditing

Auditing with zkLedger



zkLedger proposed a **privacy-preserving ledger** system to enable confidential transactions while still **allowing for auditability**.

Privacy-preserving Ledger

Asset	From	To	Amount
\$			
\$			
\$			

Auditors can still examine the **Privacy-preserving ledger** to verify to following financial invariants:



- ✓ Consent to transfer
- ✓ Has assets to transfer
- ✓ Assets neither created nor destroyed

[Motivation]

zkLedger currently only exists as an **abstract library** without a working prototype fully implemented on a blockchain platform.

The practicability of the **implementation of zkLedger on a blockchain platform with consensus** mechanism has yet to be evaluated.

Thesis Summary



This thesis has presented a comprehensive study on the design and implementation of an integrated system for privacy-preserving, and auditable decentralized transaction system on Hyperledger Fabric [Androulaki et al. 18].

Objective



- i.) Permissioned blockchain for inter-bank transactions will reveal sensitive information.
- ii.) zkLedger has yet to be implemented on a blockchain platform with consensus mechanism.

Solution



- i.) We have implemented zkLedger on top of Hyperledger Fabric, on a set of AWS EC2 VMs.
- ii.) The two approaches are proposed as trade-offs between security and efficiency:
 - Approach 1 - complete security properties but requires serial transaction processing.
 - Approach 2 - concurrent transactions with a delayed proof.
- iii.) The performance of the integrated system such as the throughput and the latency of the two approaches was evaluated.

Output



- i.) Performance would degrade as the number of participants increases in the network for both approaches.
- ii.) The second approach allows concurrent transactions, the system would more practical with a higher throughput.

Technical Background

Hyperledger Fabric

HYPERLEDGER FABRIC BLOCKCHAIN

Hyperledger Fabric [Androulaki et al. 18] is a widely adopted enterprise-grade [permissioned blockchain](#) platform, developed under the Linux Foundation's Hyperledger project. It adopts [Practical Byzantine Fault Tolerance \(PBFT\)](#) where the nodes in the system will act on the [consensus reached by the majority](#).

1. Submit Transaction

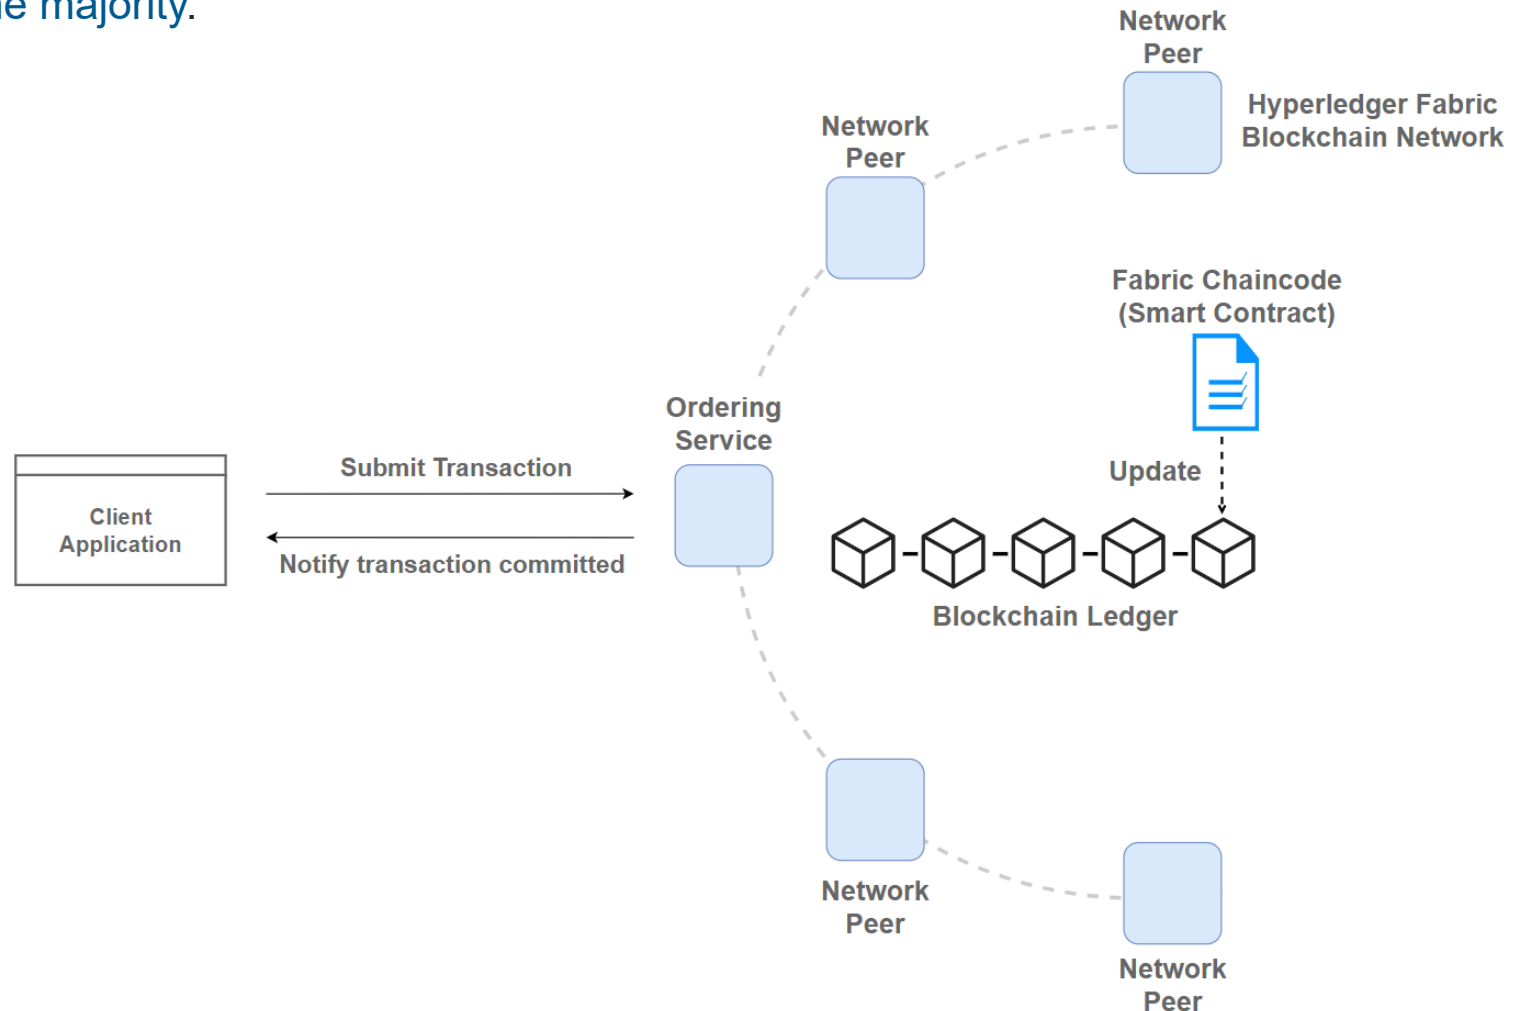
An application submits a transaction to the blockchain network.

2. Create a Block

Once the orderers have created a block, they broadcast it to the network for validation by the peers.

3. Validate Transactions

Each peer validates the transactions and appends the block to its local copy of the ledger.



Hyperledger Fabric

HYPERLEDGER FABRIC BLOCKCHAIN

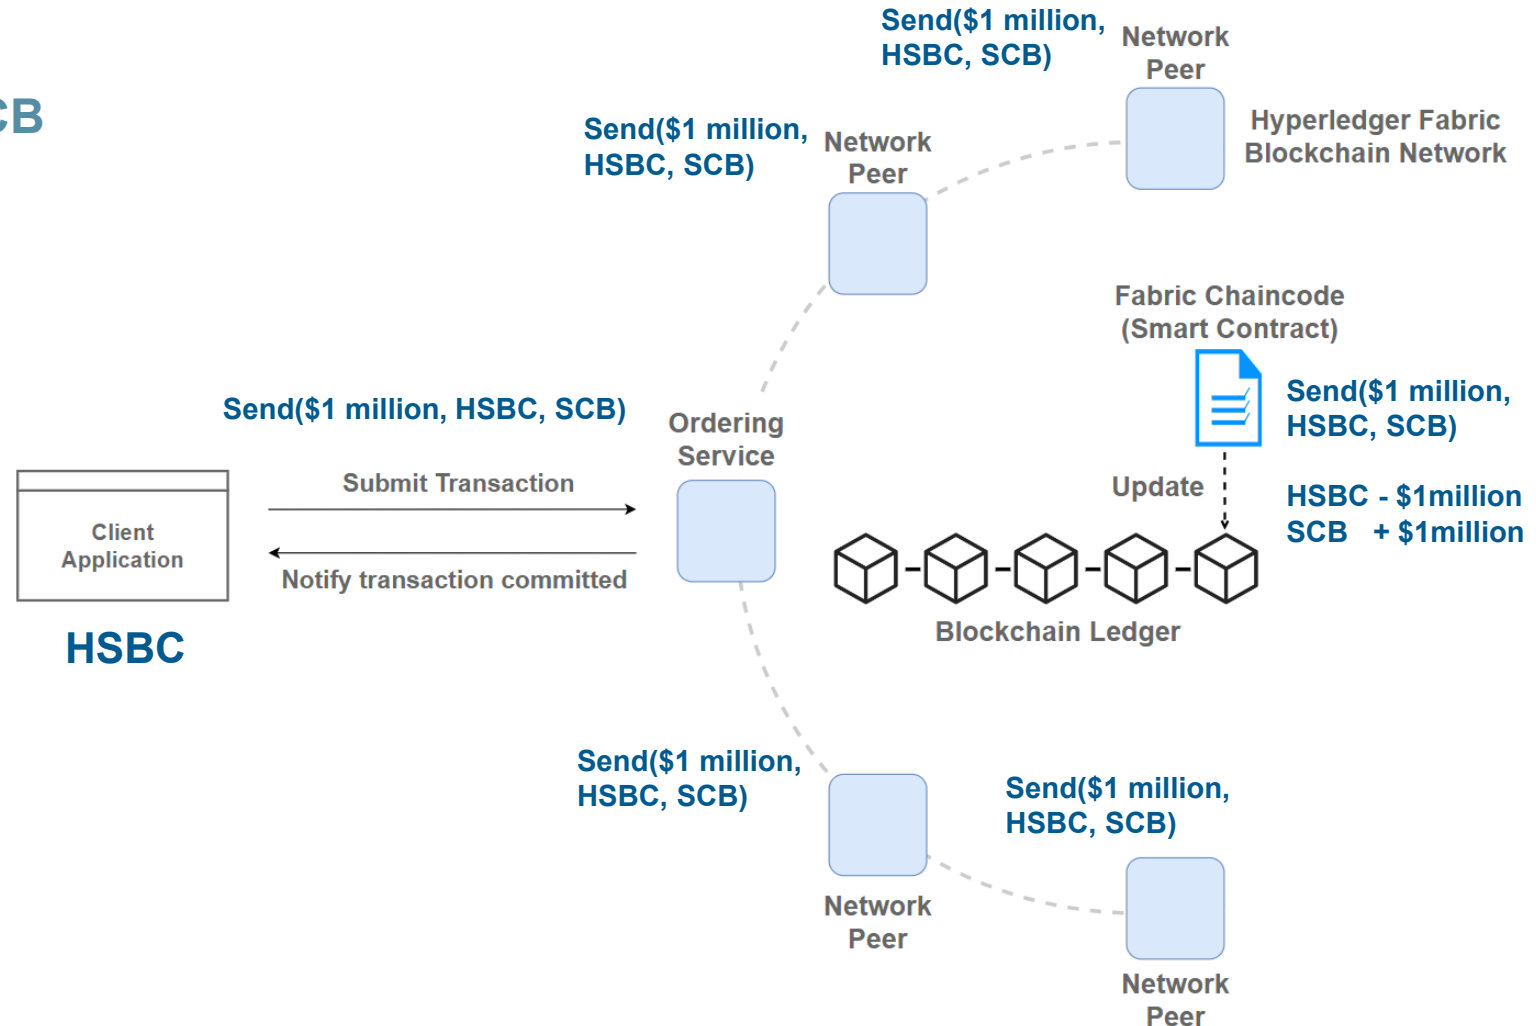
However, the **transaction content is clear** to the participants in the network, meaning all parties have complete transparency on the network.

1. HSBC wants to send \$1 million to SCB



2. HSBC submits the request (transaction) to Hyperledger Fabric

Send(\$1 million, HSBC, SCB)



zkLedger

zkLedger [Narula et al. 18] presents the first system for distributed ledgers to support **privacy-preserving transactions** with **Non-Interactive Zero-Knowledge (“NIZK”) proofs**, and **auditability**.

[Commitment Schemes]

Sender creates the Pedersen commitments:

$$cm(v) = g^v h^r$$

v	value
r	randomness
g	generator
h	generator

Hiding & Binding

The **value** committed to is **hidden** from anyone who does not know the random number used in the commitment and computationally **infeasible to change**.

Homomorphically combined

Additively homomorphic, commitments to values and randomness can create a commitment product.

$$cm_1 * cm_2 = g^{v_1} h^{r_1} * g^{v_2} h^{r_2}$$
$$cm_{1+2} = g^{v_1+v_2} * h^{r_1+r_2}$$

[Tabular Structure]

Asset	SCB	HSBC	BOC
\$	cm(-1M)	cm(1M)	cm(0)
\$	cm(2M)	cm(0)	cm(-2M)
\$	cm(-1M)	cm(1M)	cm(0)

Sender creates the latest row (SCB) →



Auditor

How many dollars do you hold?



2 million



HSBC

Open cm(1M) x cm(0) x cm(1M) to 2 million

zkLedger

zkLedger enables **public verifiability** where the general public can mathematically verify the integrity of each transaction recorded in the ledger by using **Non-Interactive Zero-Knowledge (“NIZK”) proofs**. NIZK proofs provide Verifiability without requiring interaction between the prover and verifier.

[Non-Interactive Zero-Knowledge Proofs]

Zero-knowledge proof allows one to **prove** knowledge of a secret **without revealing** any information about the secret itself.



Proof of Assets π^A

This proof ensures no party could create a transaction without owning **enough assets** to transfer.



Proof of Balance π^B

This proof ensures that a single transaction in a row conserves the overall balance in the system, meaning there is **no increase or decrease** in the overall assets in the system.



Proof of Consistency π^C

This proof ensures that the spending organization provides the **correct randomness** used in the commitment to the receiving organization.



[Proof of Assets] π^A

SCB	HSBC	BOC
cm(-1M)	cm(1M)	cm(0)
cm (2M)	cm(0)	cm(-2M)
cm(-1M)	cm(1M)	cm(0)

The **sum of the assets** in the column should be **greater than 0**.

[Proof of Balance] π^B

SCB	HSBC	BOC
cm(-1M)	cm(1M)	cm(0)
cm(2M)	cm(0)	cm(-2M)
cm(-1M)	cm(1M)	cm(0)

The **sum of the commitments** of the row should return to a commitment to a value of **0**.

[Proof of Consistency] π^C

SCB	HSBC	BOC
cm(-1M)	cm(1M)	cm(0)
cm(2M)	cm(0)	cm(-20M)
cm(-1M, r_1) Token ₁	cm(1M, r_2) Token ₂	cm(0, r_3) Token ₃

The **randomness** used in cm and the audit Token should be the same.

$$Token = pk^r$$

[Final Transaction Construction]

Asset		HSBC	SCB	BOC	...	Citibank
\$		cm(−10 ⁶ , r_1), Token ₁ , $\pi_1^A, \pi_1^B, \pi_1^C$	cm(10 ⁶ , r_2), Token ₂ , $\pi_2^A, \pi_2^B, \pi_2^C$	cm(0, r_3), Token ₃ , $\pi_3^A, \pi_3^B, \pi_3^C$...	cm(0, r_n), Token _n , $\pi_n^A, \pi_n^B, \pi_n^C$

Challenges

Challenge from Hyperledger Fabric

Multiversion concurrency control (“MVCC”) is a **concurrency control mechanism** used in Hyperledger Fabric to manage concurrent access to shared data. It allows multiple transactions to access the same data **simultaneously** without conflicts, it makes sure **no double spending or inconsistency in data** will occur.

Concurrency Issue: If the transactions of the same bank are created in the same block

HSBC wants to send \$2 million to SCB



SCB wants to send \$1 million to BOC



Initial State:

HSBC : 10 M
BOC : 10 M
SCB : 10 M

**Validate
against
initial state**

[Same Block]

HSBC : 10 M
BOC : 10 M
SCB : 10 M



**Tx 1 can
proceed to
execute**



HSBC : 8 M
BOC : 10 M
SCB : 12 M



**Tx 2 cannot
be executed**



Challenge: cannot **read / write** on data that have changed in the **same block**.

Challenge from zkLedger

The commitments are homomorphically combined to obtain a commitment to the value of the total assets. Therefore, transactions must be created **serially** to ensure that the sum of the assets is correct when producing the **Proof of Assets**.

Concurrency Issue: If the banks are creating transactions almost simultaneously

HSBC wants to send \$2 million to SCB



HSBC

SCB

SCB wants to send \$1 million to BOC



SCB

BOC

HSBC's
Client
Application

Request Tx 1

SCB's Client
Application

Request Tx 2

[Committed Transactions]

SCB	HSBC	BOC
cm(2M)	cm(0)	cm(-2M)
cm(-3M)	cm(3M)	cm(0)

[Upcoming Transactions]

cm(2M)	cm(-2M)	cm(0)	✓
"Able to create the correct Proof of Assets π^A since I know all the previous transactions"			
cm(-1M)	cm(0)	cm(1M)	✗

"**NOT** able create the correct Proof of Assets π^A since I do **NOT** know all the previous transactions"

Challenge: cannot create a correct Proof of Asset π^A if a sender has just received a transaction.

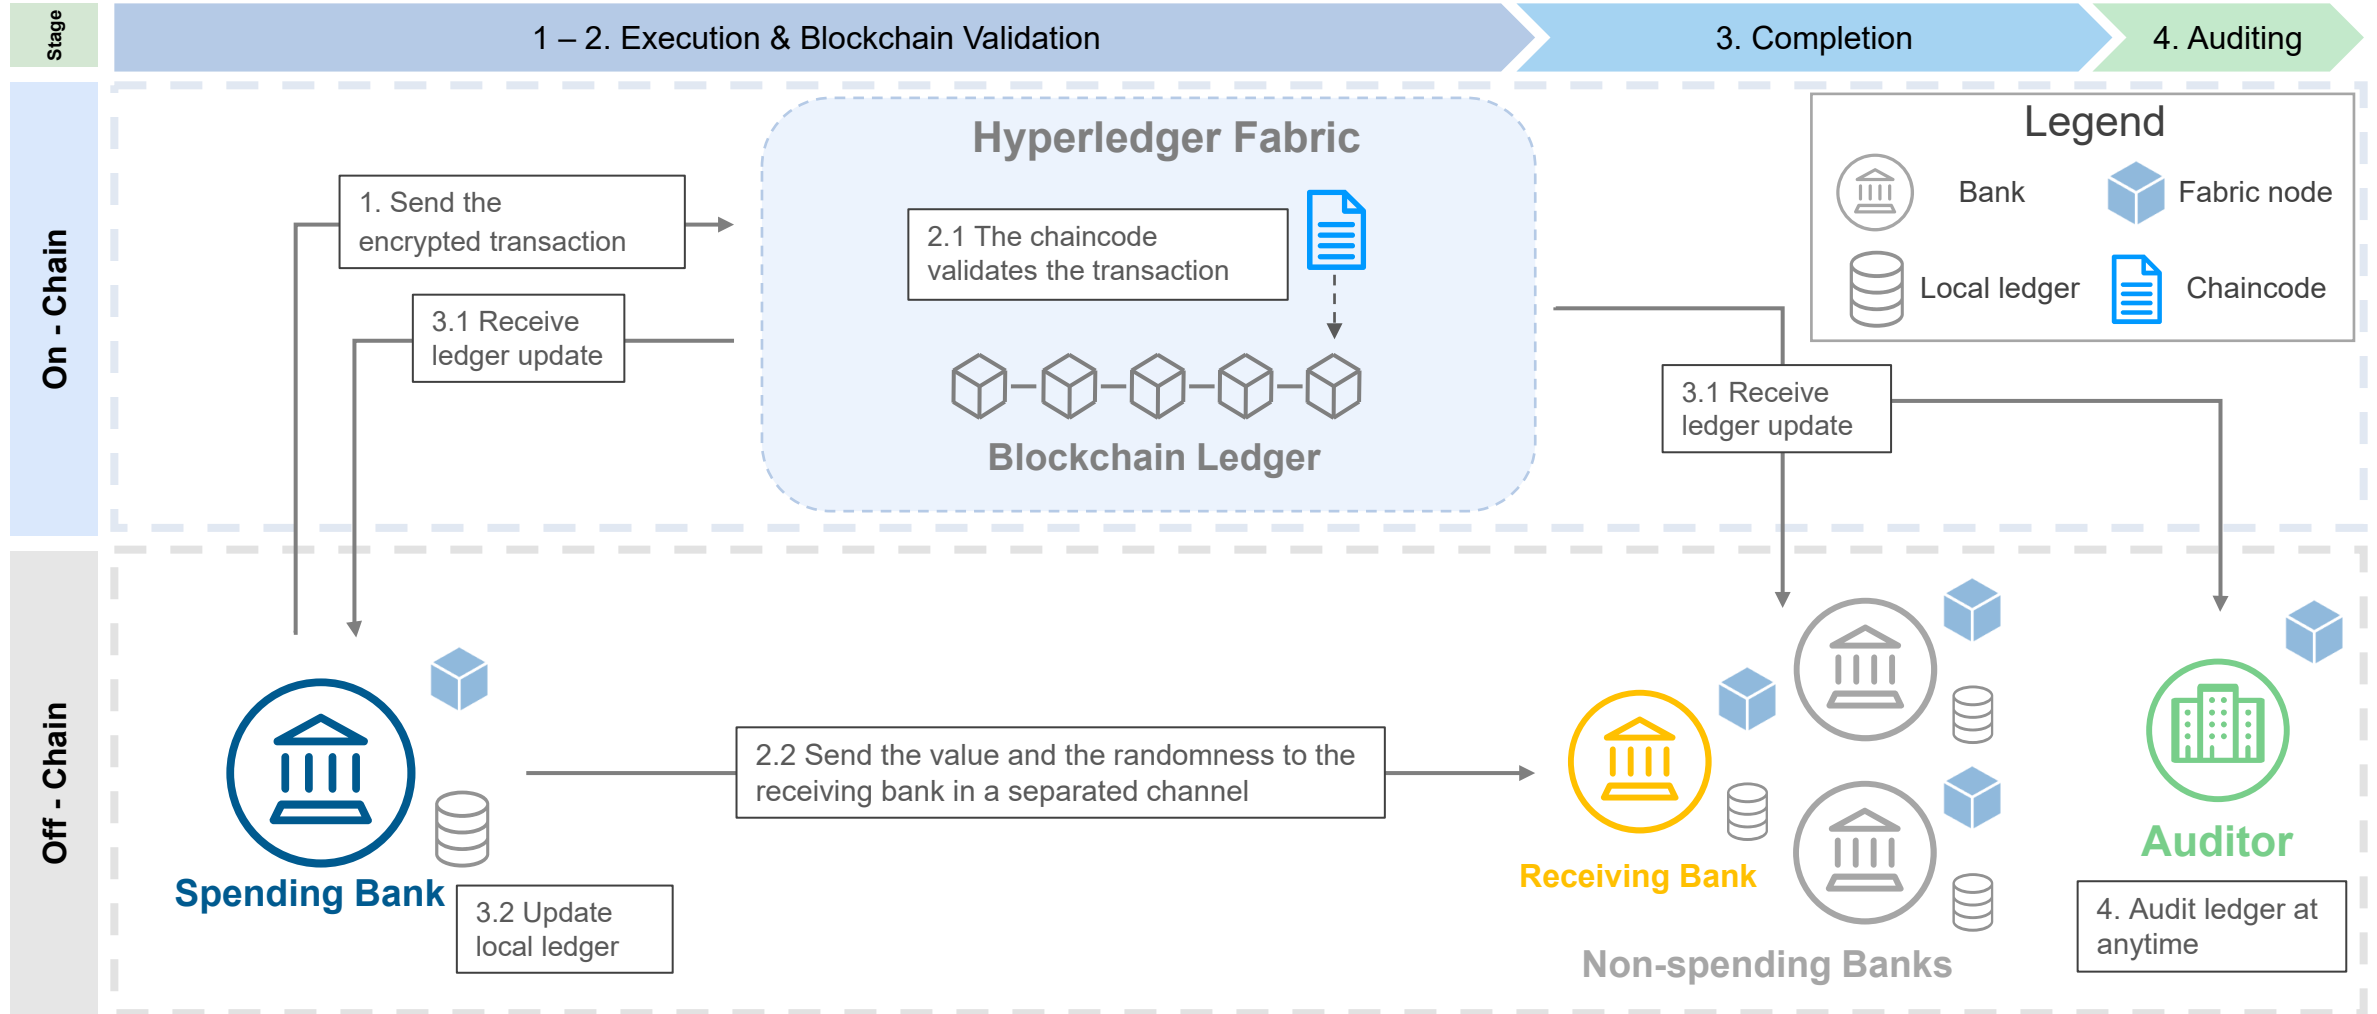


Our Solution:

An integrated system for privacy-preserving, and auditable transactions on Hyperledger Fabric

System Design

Our design features a secure transaction system that operates on a [Hyperledger Fabric](#) blockchain network with bank clients. Each bank's machine has a node, and a client application leveraging the [zkLedger](#) cryptographic library.



Objective

zkLedger currently only exists as an **abstract library** without a working prototype of the consensus mechanism included. Additionally, together with Hyperledger Fabric, it has **concurrency issues** that need to be addressed for it to be a viable solution. Hence, we are proposing two approaches to address the challenges.

Approach 1 - Serial Transactions

This approach involves implementing zkLedger on Hyperledger Fabric **without altering the security properties**, resulting in a **fully secure but less scalable** system due to the serial transaction processing requirement.

This approach may be suitable for applications where **security is of paramount importance**, and transaction throughput is not a critical concern.

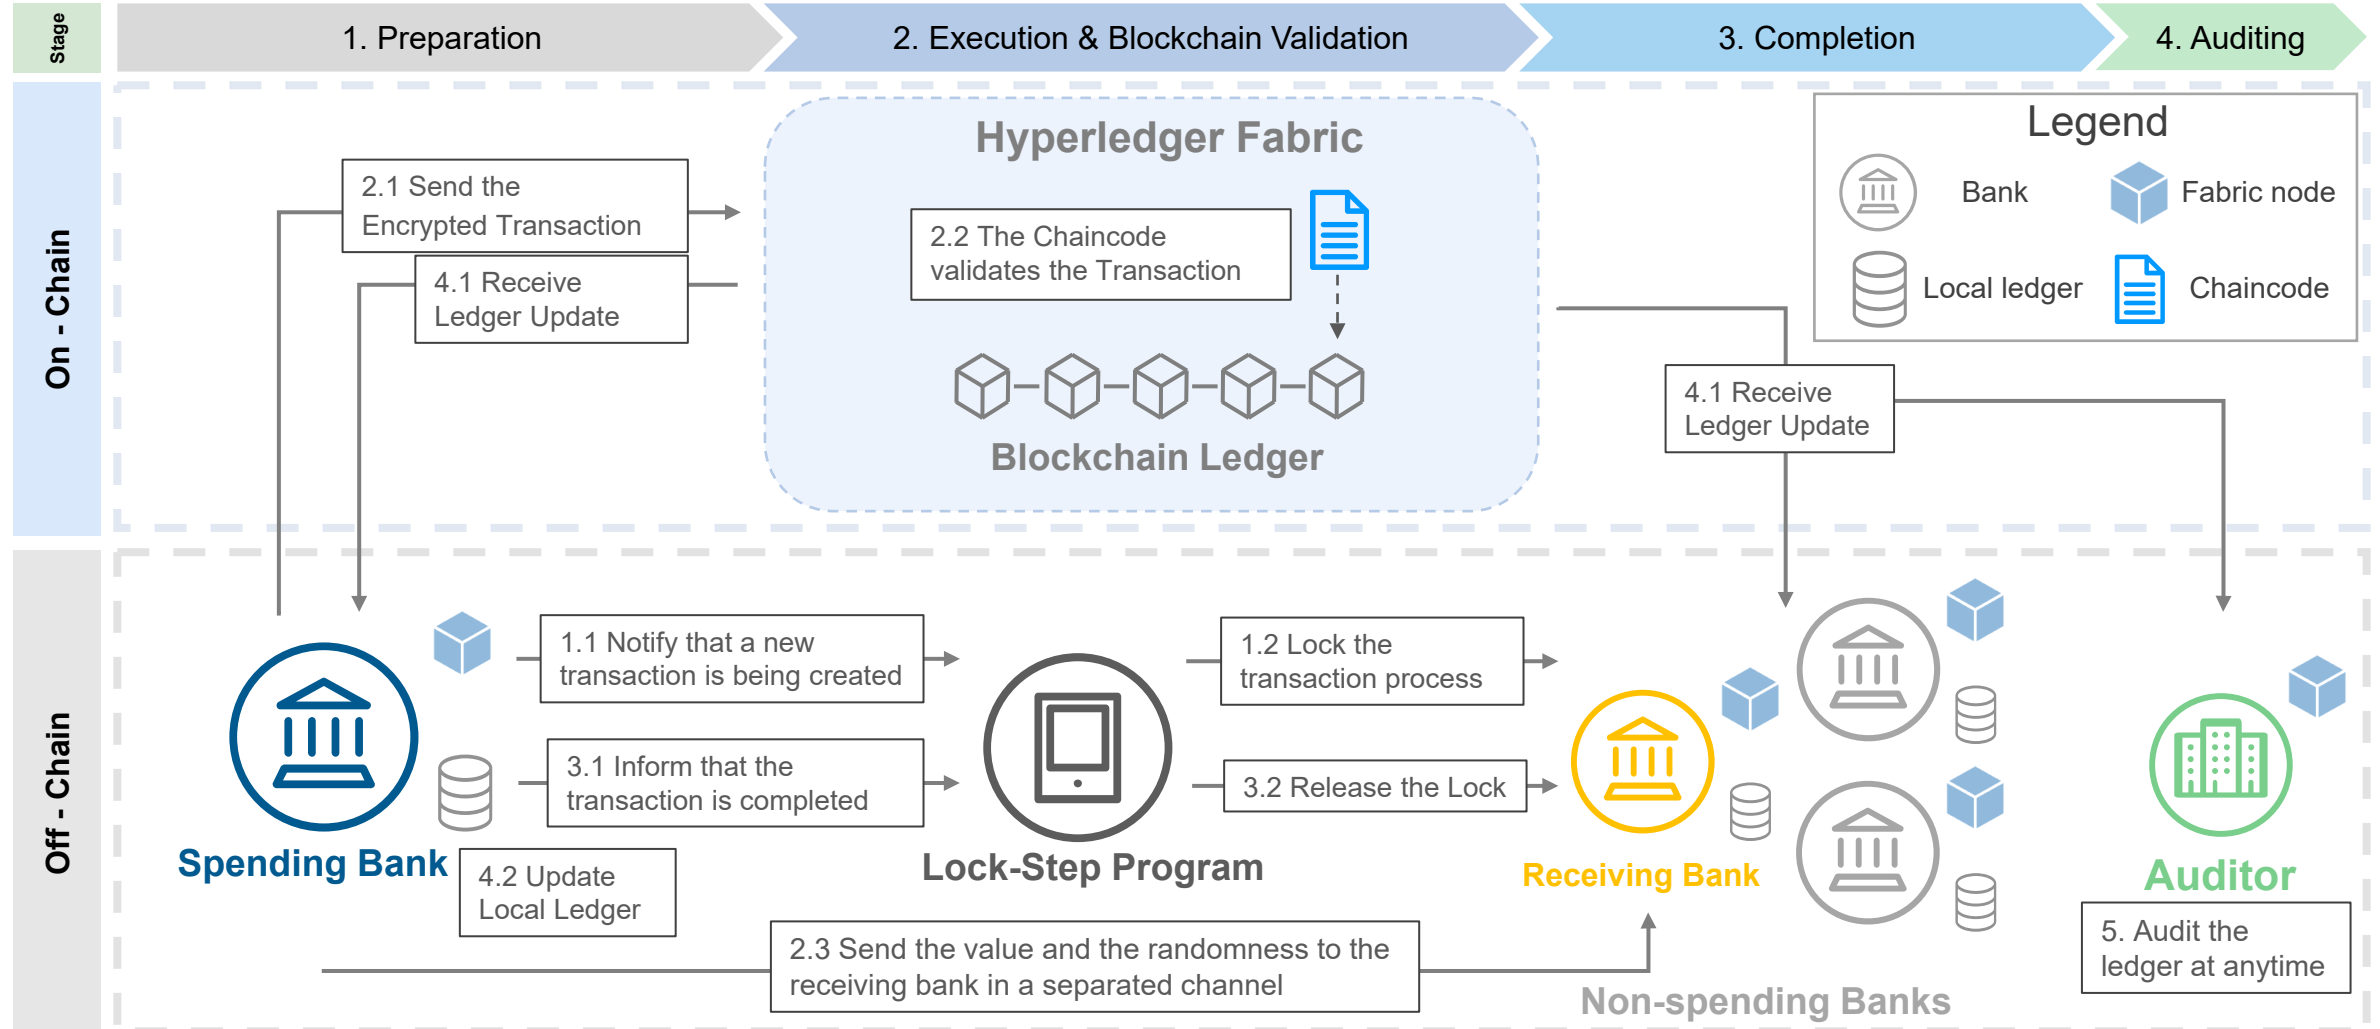
Approach 2 - Full Concurrency

This approach introduces a **delayed verification** of one of the zkLedger security properties, **Proof of Assets**, allowing participants to **create transactions concurrently**.

This design significantly **improves the performance** of the system, as participants no longer need to wait for the completion of the previous transaction before initiating new ones.


Approach 1 - Serial Transactions

This approach enables the system to attain **complete verifiability** by validating all NIZK proofs. A lock-step program is established to facilitate communication among banks, ensuring the accurate sequencing of transactions.



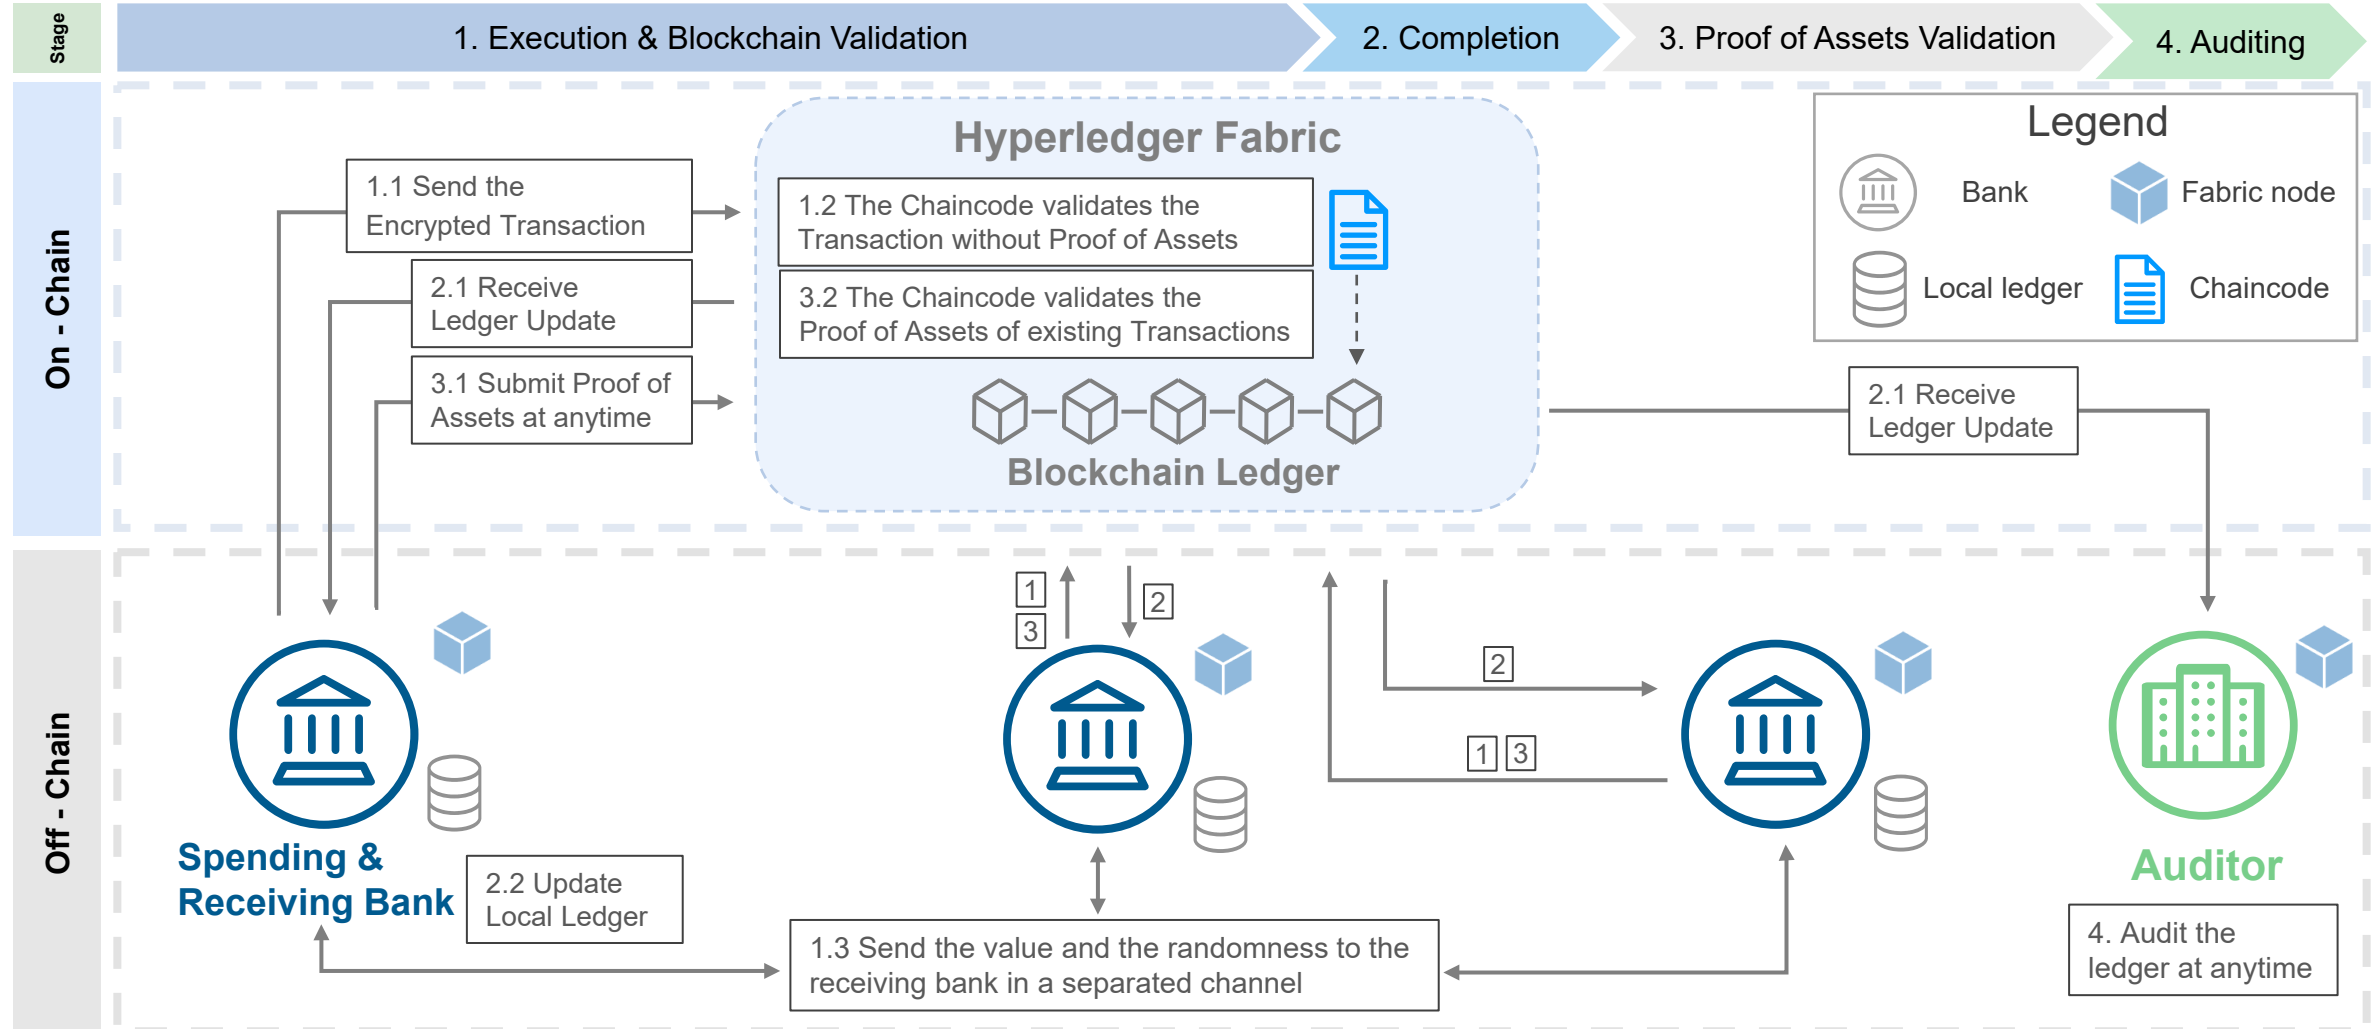
Approach 1 - Serial Transactions

A high-level block data structure for the blockchain in the [Serial Transactions](#) approach. The block data consists of a Fabric transaction, which includes the banks' commitment caches and audit token caches, and [one](#) zkLedger transaction containing [all the proofs](#).




Approach 2 - Full Concurrency

This approach enables the system to attain **full concurrency** by delaying the on-chain validation of Proof of Assets π^A due to its serialized characteristic. The transaction structure remains the same, ensuring provability and publicly verifiability.



Approach 2 - Full Concurrency

A high-level block data structure for the blockchain in the Full Concurrency approach. The block data includes a Fabric transaction containing **multiple** zkLedger transactions **without Proof of Assets** π^A or an update of the **previous** zkLedger transaction **with Proof of Assets** π^A attached.



Approaches Comparison



The two approaches has mainly demonstrated the **trade-offs** below in the design and implementation of such a privacy-preserving transaction system.

Approach 1 - Serial Transactions

Approach 2 - Full Concurrency

[Comparison]

i. Full Verifiability

The chaincode can **validate** all of the proofs.

ii. Less Efficiency

Banks must create transactions **serially**.

iii. Less resource consumption

As the blockchain only process **one transaction per block**.

iv. Straightforward implementation

The **cryptographic components** were not modified much from the zkLedger's cryptographic library.

i. Delayed Verifiability

Proof of Assets, is **delayed for verification**.

ii. Enhanced Efficiency

Participants can **create transactions concurrently**.

iii. High resource consumption

The blockchain might process **many transactions in one block**.

iv. Complex implementation

A **great amount of modification was made** from the zkLedger's cryptographic library.

Evaluation

System's Performance Evaluation



In our experimental setup, we have designed and implemented a working prototype of the proposed application over [Hyperledger Fabric v1.4](#) as the underlying blockchain platform, deployed on a set of AWS EC2 instances.

System: AWS EC2 c5.4xlarge instance type, 16 vCPUs with 3.6GHz clock speed, 32GB of RAM

OS: 64-bit Linux 5.4.0-1051-aws kernel with Ubuntu 18.04.1 as the operating system

Crypto libraries: zkSigma, btcec, secp256k1

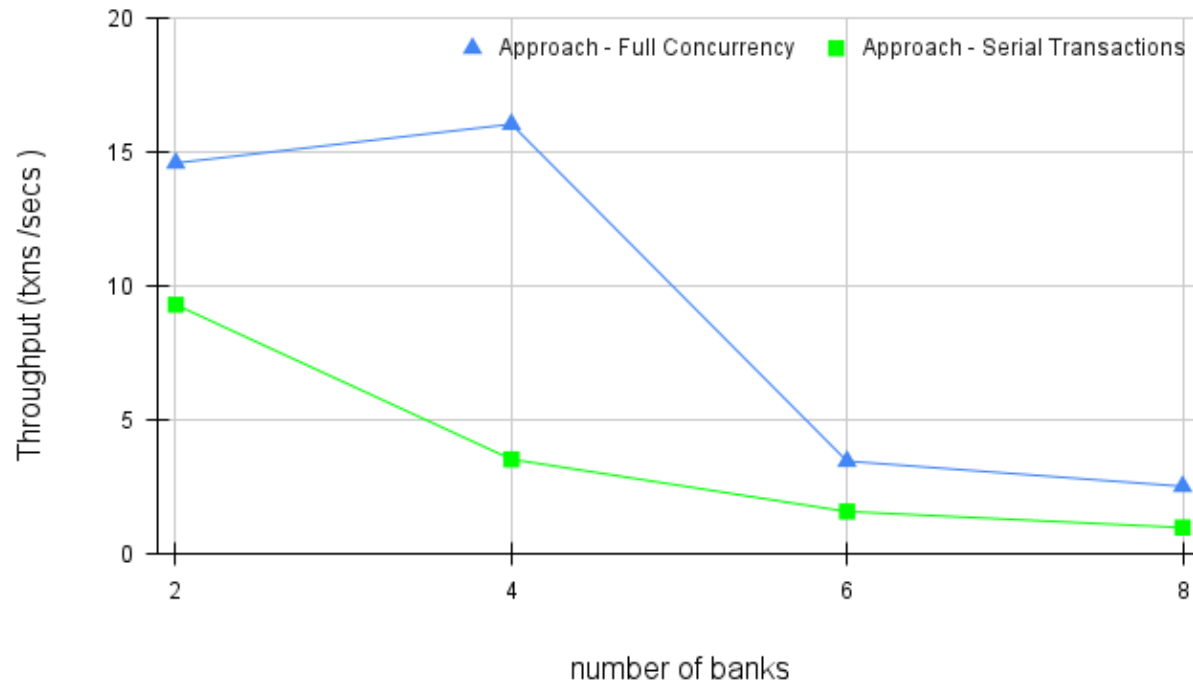
Client application and chaincode: Go version 1.9

System's Performance Evaluation

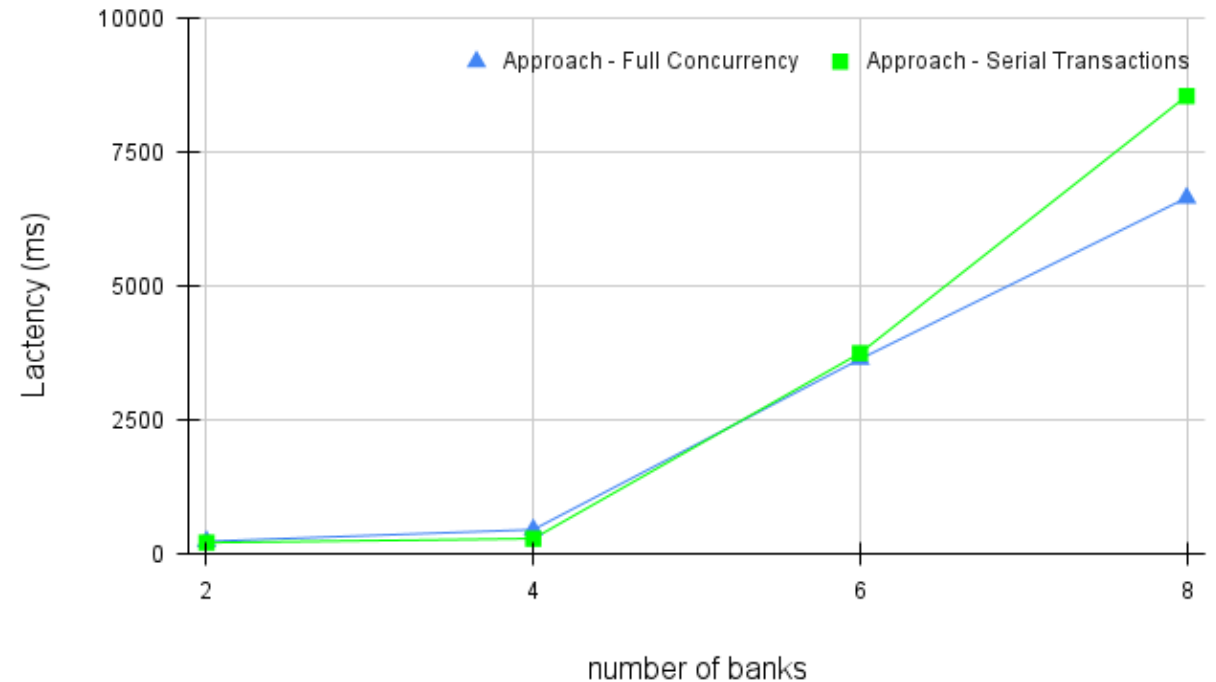
The below graphs show the **throughput** and the **end-to-end latency** measured in both the Serial Transactions and Full Concurrency approaches.

In the **Serial Transactions** approach, the banks send **one transaction at a time** across participants.

In the **Full Concurrency** approach, the banks send **N transactions simultaneously** where N refers to the number of participating banks.



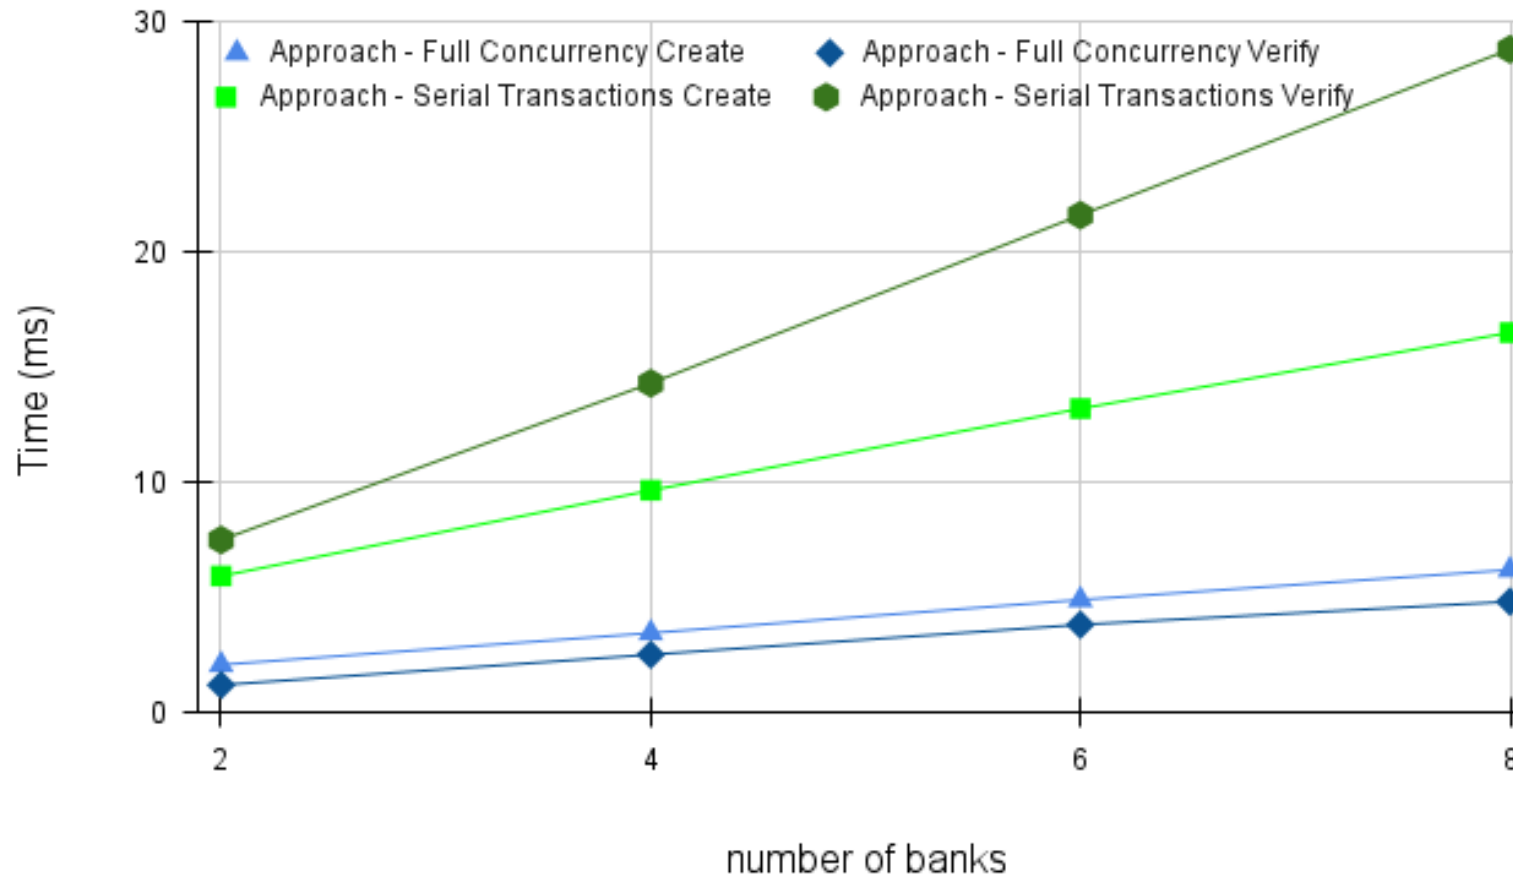
[Throughput evaluation vs number of banks]



[Latency evaluation vs number of banks]

System's Performance Evaluation


The below graph shows the time it takes for the banks to **create** the transactions, and the time it takes for the Hyperledger Fabric to **validate** the transactions in both the Serial Transactions and Full Concurrency approaches.



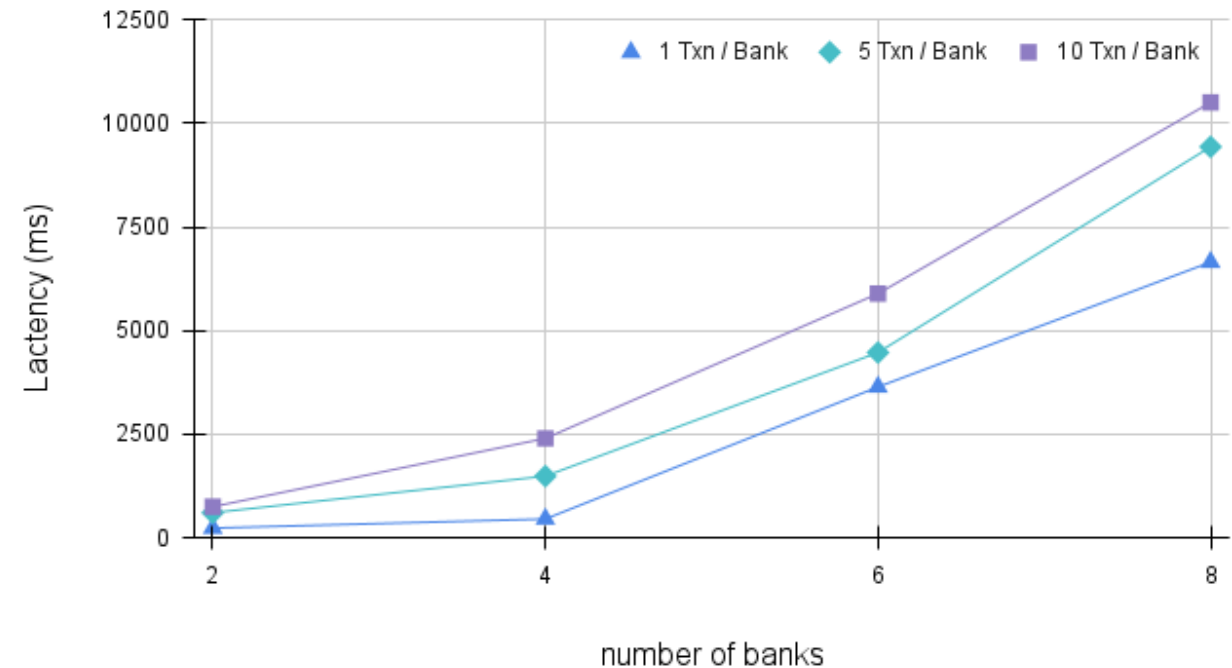
[Transaction create and verify time evaluation vs number of banks]

System's Performance Evaluation

Since banks have the ability to submit as many transactions as they wish in the Full Concurrency approach, it would be interesting to see how the system would react depending on the number of transactions each bank is transacting simultaneously.



[Throughput evaluation vs number of banks]



[Latency evaluation vs number of banks]

Conclusion

Conclusion



This thesis has presented a comprehensive study on the design and implementation of **an integrated system for privacy-preserving, and auditable decentralized transaction system on Hyperledger Fabric with zkLedger.**

Our result

- We have **implemented zkLedger on top of Hyperledger Fabric** on a set of AWS EC2 machines to **evaluate the system performance** as a distributed inter-bank transaction system.
- The limitation of concurrency was identified, we have proposed of **two distinct approaches** to address the design practicality and the identified limitations.

Core finding

- The two approaches are proposed as **trade-offs between security and efficiency**: the first provides complete security properties, while the second offers increased system performance.
- Although the second approach has increased system performance, in general, the **system scalability has to be further improved** to optimize the performance in a larger group of participants.

Future Direction

- Explore different cryptographic approaches of **privacy and auditability** suitable for inter-bank transaction system.



Q&A




This page intentionally left blank.

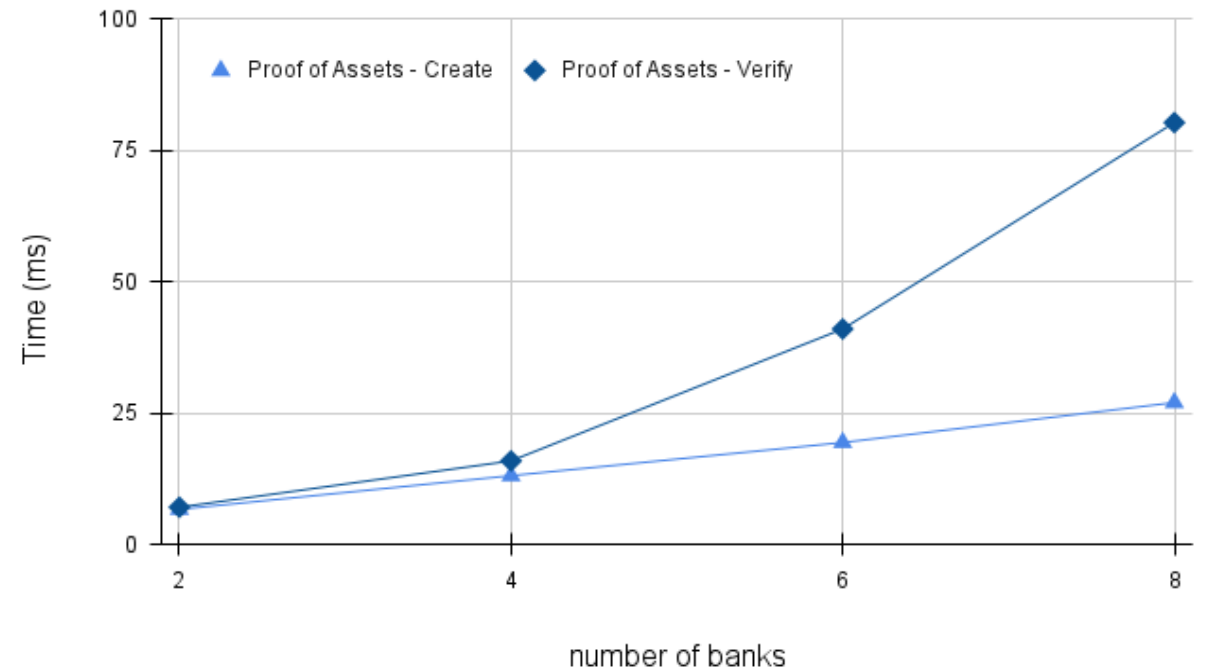
| Appendix

System's Performance Evaluation

The below graphs show the **throughput** measured for the delayed **Proof of Assets transactions** of the Full Concurrency approach depending on the number of banks. Each bank is set to create and submit **10 PoA** transactions as one batch.



[PoA Throughput evaluation vs number of banks]



[PoA Transaction create and verify time evaluation vs number of banks]